

# 1. REINFORCE

**Summary:** REINFORCE introduced the idea of directly optimizing a policy using the policy gradient theorem. This method represented a significant shift from value-based methods by focusing on optimizing the policy itself.

**Big Idea:** The major change was moving away from value function-centric methods (like Q-learning) to directly optimizing the policy via gradient ascent. This allowed for more complex and flexible policies.

**Motivation:** REINFORCE addressed the difficulty of learning optimal policies directly, especially in environments with large or continuous state-action spaces.

**Appropriateness:** Suitable for problems with large state spaces or continuous action spaces where a direct parameterization of the policy is feasible.

**Handling States, Time, Actions, Rewards:** It parameterizes the policy and updates it based on rewards received. The return  $R_t$  from the start of an episode is used to evaluate the policy's performance.

**Mathematical Formulation:**

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) R_t \right]$$

The gradient  $\nabla_{\theta} J(\theta)$  reflects how policy parameters  $\theta$  should be adjusted to maximize the expected return.

## 2. DQN, DDQN, Prioritized Replay

### DDQN

**Summary:** DQN (Deep Q-Network) uses a neural network to approximate Q-values, addressing the instability in Q-learning with function approximation. DDQN (Double DQN) mitigates Q-value overestimation by decoupling action selection and evaluation. Prioritized Replay DDQN improves efficiency by prioritizing important experiences.

**Big Idea:** The key innovation was the use of neural networks to approximate Q-values and the introduction of experience replay and target networks to stabilize training.

**Motivation:** These methods were developed to handle high-dimensional state spaces and stabilize learning, which earlier algorithms struggled with.

**Appropriateness:** Effective for problems with discrete action spaces and large state spaces, such as video games or complex control tasks.

**Handling States, Time, Actions, Rewards:** DQN uses a Q-network to approximate action-value functions. DDQN improves this by using a separate target network for stability. Prioritized Replay enhances efficiency by sampling experiences based on their importance.

**Mathematical Formulation:**

$$L(\theta) = \mathbb{E}_{(s,a,r,s')} \left[ \left( r + \gamma \max_{a'} Q_{\theta-}(s', a') - Q_{\theta}(s, a) \right)^2 \right]$$

The loss function  $L(\theta)$  measures the difference between the predicted Q-value and the target Q-value, driving the update of the Q-network parameters  $\theta$ .

## 3. A2C/A3C

**Summary:** A2C (Advantage Actor-Critic) uses both an actor to update the policy and a critic to evaluate it. A3C (Asynchronous Actor-Critic Agents) extends A2C by running multiple agents in parallel to improve stability and learning speed.

**Big Idea:** The innovation lies in the use of both actor and critic networks to separately handle policy improvement and value estimation, and the parallelization in A3C to stabilize learning.

**Motivation:** These methods were designed to overcome the high variance in policy gradient methods and to speed up training through parallelism.

**Appropriateness:** Suitable for environments where multiple agents can interact concurrently, and where the learning process can benefit from both policy and value function optimization.

**Handling States, Time, Actions, Rewards:** A2C uses the advantage function  $A(s_t, a_t)$  to guide policy updates and evaluates state values through the critic. A3C utilizes multiple agents to collect diverse experiences.

**Mathematical Formulation:**

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{s_t, a_t} [\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) A(s_t, a_t)]$$

$$L(\phi) = \mathbb{E}_{s_t} \left[ \left( R_t - V_{\phi}(s_t) \right)^2 \right]$$

The actor update uses the advantage function to guide policy improvement, while the critic update minimizes the error between predicted and actual returns.

## 4. DPG, DDPG

**Summary:** DPG (Deterministic Policy Gradient) and DDPG (Deep Deterministic Policy Gradient) extend Q-learning to continuous action spaces. DDPG enhances DPG with experience replay and target networks for improved stability.

**Big Idea:** The key innovation is the use of deterministic policy gradients for continuous action spaces and the introduction of experience replay and target networks in DDPG.

**Motivation:** These methods were developed to handle continuous action spaces and to stabilize training through improved learning techniques.

**Appropriateness:** Suitable for problems with continuous action spaces, such as robotic

control or autonomous driving.

**Handling States, Time, Actions, Rewards:** DDPG uses actor-critic architecture with continuous action spaces, leveraging target networks and experience replay to stabilize learning.

**Mathematical Formulation:**

$$\nabla_{\phi} J(\phi) = \mathbb{E}_s \left[ \nabla_a Q_{\theta}(s, a) \Big|_{a=\mu_{\phi}(s)} \nabla_{\phi} \mu_{\phi}(s) \right]$$

$$L(\theta) = \mathbb{E}_{(s,a,r,s')} \left[ \left( r + \gamma Q_{\theta-}(s', \mu_{\phi-}(s')) - Q_{\theta}(s, a) \right)^2 \right]$$

The actor update uses deterministic policy gradients, while the critic loss function updates the Q-network based on target values from the target networks.

## 5. TRPO, PPO

**Summary:** TRPO (Trust Region Policy Optimization) ensures policy updates stay within a trust region to improve stability, while PPO (Proximal Policy Optimization) simplifies TRPO with a clipped objective to constrain updates.

**Big Idea:** TRPO's innovation was to use a trust region constraint for policy updates, while PPO simplified this with a clipping mechanism, making it more practical and easier to implement.

**Motivation:** Both methods were developed to address issues of instability and inefficiency in policy optimization, particularly in high-dimensional spaces.

**Appropriateness:** Effective for high-dimensional and continuous action spaces where stability and efficiency in policy optimization are crucial.

**Handling States, Time, Actions, Rewards:** Both methods use the policy gradient approach to optimize policies. PPO simplifies TRPO by using a clipped objective function to limit policy changes.

**Mathematical Formulation:**

$$\max_{\theta} \mathbb{E}_{s,a \sim \pi_{\theta_{\text{old}}}} \left[ \frac{\pi_{\theta}(a|s)}{\pi_{\theta_{\text{old}}}(a|s)} A^{\pi_{\theta_{\text{old}}}(s,a)} \right]$$

$$\max_{\theta} \mathbb{E}_{s,a \sim \pi_{\theta_{\text{old}}}} [\min(r_t(\theta) A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) A_t)]$$

TRPO uses a trust region constraint to prevent large policy updates, while PPO uses a clipped objective function to achieve similar stability.

## 6. MCTS

**Summary:** Monte Carlo Tree Search (MCTS) is a search algorithm for making decisions in large spaces using simulation to evaluate potential actions.

**Big Idea:** The fundamental innovation is the use of simulation to explore and evaluate decisions, enabling effective decision-making in complex and large state spaces.

**Motivation:** MCTS was developed to make decisions in domains with large state spaces where traditional methods would be computationally prohibitive.

**Appropriateness:** Suitable for decision-making problems with large state spaces and where simulations can be used to evaluate potential actions, such as in game-playing and planning problems.

**Handling States, Time, Actions, Rewards:** MCTS builds a search tree where nodes represent states and edges represent actions. It uses simulations to estimate the value of different actions.

**Mathematical Formulation:**

$$Q(s, a) = \frac{w(s, a)}{n(s, a)} + c \sqrt{\frac{\ln N(s)}{n(s, a)}}$$

The formula balances exploitation (average reward) and exploration (uncertainty) using the UCT (Upper Confidence Bound for Trees) algorithm.

## 7. AlphaGo, AlphaZero

**Summary:** AlphaGo and AlphaZero use deep learning combined with MCTS for game-playing. AlphaGo was designed for Go, while AlphaZero generalizes this approach for various games using self-play.

**Big Idea:** The innovation is combining deep learning with MCTS to create powerful game-playing agents. AlphaZero extended this to general game-playing with self-play.

**Motivation:** These methods address the need for high-performance game-playing agents that can learn and improve from self-play, avoiding reliance on human expertise.

**Appropriateness:** Ideal for strategic games and other complex environments where self-play and simulation are feasible for training.

**Handling States, Time, Actions, Rewards:** AlphaGo and AlphaZero use deep networks to estimate policy and value functions, combined with MCTS for decision-making.

**Mathematical Formulation:**

$$L(\theta) = \mathbb{E}_{(s,\pi,z) \sim D} \left[ (z - V(s; \theta))^2 - \pi \cdot \log \pi(a|s; \theta) + \lambda \|\theta\|^2 \right]$$

The loss function balances prediction of state values, policy probabilities, and regularization, guiding the training of the policy and value networks.